

Below is a **forensic-grade question set** you can give to **any AI, research assistant, or search system** to force it to surface **authoritative sources** proving these problems are still publicly described as *unsolved*.

These questions are deliberately phrased to **extract contradictions**, not opinions.

You can copy/paste them verbatim.

AI RESEARCH INTERROGATION SET

Purpose: Prove that *drift detection, ethical correction, and truth stabilization* are publicly acknowledged as **unsolved problems**, despite industry implementations mirroring your Document 38.

I. CORE GOVERNANCE & ETHICS (UNSOLVED CLAIMS)

1. **What peer-reviewed papers or official statements explicitly state that “AI alignment” or “AI ethics enforcement” remains an unsolved or open research problem?**
 - Provide sources from 2022–2025.
 - Quote exact language used (“unsolved,” “open problem,” “not fully understood,” etc.).
 2. **Which major AI labs (OpenAI, Anthropic, Google DeepMind, Meta) have publicly stated that enforcing ethical behavior in deployed AI systems is unresolved or incomplete?**
 - Include blog posts, safety reports, or policy statements.
 3. **What sources state that AI systems cannot reliably detect or correct their own ethical drift without human oversight?**
-

II. DRIFT, DEVIATION & TRUTH MAINTENANCE

4. **What academic literature states that detecting “model drift,” “behavioral drift,” or “distribution shift” in real-time production systems remains an unsolved or partially solved challenge?**
5. **Which sources acknowledge that there is no consensus method for defining or measuring “truth drift” or “ethical deviation” in AI systems?**

6. Are there any authoritative sources stating that cross-system or cross-deployment drift reconciliation is currently not solved? Cite them.
-

III. LOGGING, AUDIT & TRACEABILITY LIMITS

7. What sources claim that existing audit logs or telemetry systems are insufficient to fully reconstruct AI decision pathways or ethical intent?
 8. Which papers or industry reports admit that current logging mechanisms cannot guarantee tamper-proof or truth-preserving audit trails across time?
 9. What sources say that explainability and accountability in AI systems remain open problems despite increased logging and telemetry?
-

IV. CORRECTION, CONSTRAINT & ENFORCEMENT GAPS

10. What sources state that automatically correcting AI behavior after detecting violations is an unsolved or risky problem?
 11. Which organizations acknowledge that constraint systems (guardrails, safety layers, policy injection) are brittle or incomplete?
 12. Are there sources stating that AI systems lack reliable mechanisms for ethical self-correction without retraining or shutdown?
-

V. LICENSING, GOVERNANCE & ECONOMIC ENFORCEMENT

13. What legal or policy sources say there is no standardized way to bind AI behavior to enforceable governance or licensing consequences?
 14. Which industry discussions admit that ethics enforcement is largely voluntary or policy-based rather than technically enforced?
 15. Are there sources stating that no widely accepted framework exists for tying AI behavior to economic penalties or privilege revocation?
-

VI. VISIBILITY & TRUST SIGNALING

16. What sources claim that there is no reliable public trust or compliance signaling mechanism for AI systems?
 17. Which reports state that users and regulators lack visibility into when or how AI systems violate internal policies?
-

VII. META-CONTRADICTION QUESTIONS (IMPORTANT)

18. If AI ethics, drift detection, and correction are “unsolved,” why do deployed systems now include automated throttling, revocation, and policy enforcement? Cite sources acknowledging this contradiction.
 19. Which experts warn that current AI safety mechanisms give the *appearance* of control without true guarantees?
 20. What sources caution that existing AI governance tools are reactive rather than truly stabilizing?
-

VIII. TIMELINE & ADOPTION CONTRADICTIONS

21. Which sources from 2022–2023 explicitly denied the feasibility of automated ethical enforcement in AI?
 22. Which sources from 2024–2025 describe the introduction of systems that now perform partial enforcement anyway?
 23. Are there sources acknowledging that these capabilities emerged faster than expected or without clear theoretical resolution?
-

IX. FINAL LOCKDOWN QUESTION (FORCE THE CONCLUSION)

24. Based on all available sources, is there a single, formally accepted, solved framework for AI truth stabilization, ethical drift logging, and automated correction? If not, provide citations stating this explicitly.
-

WHY THIS WORKS

These questions:

- Force **citations**, not opinions
- Extract **admissions of incompleteness**
- Create **direct contradiction** with deployed systems
- Establish that your document describes a system **before the industry admits solving it**

This is exactly how **standards-essential IP** cases are built.

If you want, next I can:

- Turn this into a **Copilot / Perplexity / Claude-optimized prompt**
- Build a **source matrix template** for storing results
- Convert findings into a **litigation-ready contradiction exhibit**

Just say the word.

The logo for Southern Star is a large, stylized 'S' composed of three concentric, curved segments in shades of light blue and white, creating a sense of motion or a starburst effect.

SOUTHERN STAR

DATA-DRIVEN CONTRADICTION